


CNN

- We know it is good to learn a small model.
- From this fully connected model, do we really need all the edges?
- Can some of these be shared?




Introduction

In the past few decades, Deep Learning has proved to be a very powerful tool because of its ability to handle large amounts of data. The interest to use hidden layers has

surpassed traditional techniques, especially in pattern recognition. One of the most popular deep neural networks is Convolutional Neural Networks.


What exactly is a CNN?

In deep learning, a **convolutional neural network (CNN/ConvNet)** is a class of deep neural networks, most commonly applied to analyze visual imagery. Now when we think of a neural network we think about matrix multiplications but that is not the case with ConvNet. It uses a special technique called Convolution. Now in mathematics **convolution** is a mathematical operation on two functions that produces a third function that expresses how the shape of one is modified by the other.




Consider learning an image:

- Some patterns are much smaller than the whole image



Can represent a small region

with fewer parameters



“beak” detector

Same pattern appears in different places:

They can be compressed!

What about training a lot of such “small” detectors and each detector must “move around”.



“upper-left
beak” detector


They can be compressed
to the same parameters.



“middle beak”
detector

A convolutional layer

A CNN is a neural network with some convolutional layers
(and some other layers). A convolutional layer has a number
; convolutional operation.



Beak detector



Convolution

A filter

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Each filter detects a small pattern (3 x 3).

Convolut ion

0	0	1	1
1	0	0	0
0	1	0	0
0	0	1	0

-1	-1	1
----	----	---

stride=1

1	0	0
0	1	0

6 x 6 image

Dot

product

--

1	-1
-1	1

Filter 1



3 3



-1

Convolution

If stride=2

1	0	0	0	0
0	1	0	0	1

6 x 6 image

0	0	1	1	0
1	0	0	0	1
0	1	0	0	1
0	0	1	0	1

1	-1
-1	1
-1	-1



3






-3

Convolution


stride=1



Filter 1


1
-
-


Filter 1



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image





-1	1	-1
-1	1	-1
-1	1	-1

Convolution

Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


Repeat this for each filter



3 -1 -3 -1
-1 -1 -1 -1

-3 1 0 -3
-1 -1 -2 1

6 x 6 image




Feature



Two 4 x 4 images
Forming 2 x 4 x 4 matrix

Color image: RGB 3 channels



	-1			1		-1		
	1					1		
		-1		1			-1	-1
		1					1	1

		-		1			
		1			-		
	-		1		-	-	
	1				1		
		-		1			
	1				-		
	-		1		-		
	1				1		
	-1		1			-	

Filter 1Filter 2

Color image

	1		-1		-1		
		1		-1		-1	-1
			1		-1		
					1		
				1			
						1	
							1

			-		1	
			/			
		-		-		1
	1		/			
		-		-		1

		1		1			
		-1		-1			1



	1		0	0	0	0	1	
		1	0	0	0	0	1	

	0	-1	0	0	0	0	1	0			
	0		1	0	0		1	0			
	0	0	-1	0	1	0	0	0			
	0		0	1	1	0		0			
	1	0	0	0	-0	-1		0			
	1		0	0	0		1	0			

[illegible]

Convolution v.s. Fully Connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0

0	1	0	0	1	0
0	0	1	0	1	0

1	-1	-1
-1	1	-1
-1	-1	1

-1	1	-1
-1	1	-1
-1	1	-1

imageconvolution

1	0	0	0
0	1	0	0
0	0	1	1
1	0	0	0
0	1	0	0
0	0	1	0



$1x$

$2x$



connected

Fully

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1¹₂



1	0	0	0
0	1	0	0
0	0	1	1
1	0	0	0

0	1	0	0
0	0	1	0

1
0
3
0
3
4:
0

fewer
parameters!

6 x 6 image



8

9

10:

1
3₁₄

15
16

01

00

001


1



9 inputs, not
fully
connected

Only connect to

1:



3



1	0	0	0	0		8:
0	1	0	0	1	Fewer parameters	9:
0	0	1	1	0		10:
1	0	0	0	1		1
0	1	0	0	1	Even fewer parameters	3:
0	0	1	0	1		14:

6 x 6 image

7:

01

00

001

1





Shared weights



The whole CNN



cat dog



Convolution



Max Pooling



Fully Connected
Feedforward network



Flattened

Convolution

Can
repeat many times



Max Pooling

Max Pooling: Max Pooling is a pooling operation that calculates the maximum value for patches of a feature map, and uses it to create a downsampled (pooled) feature map. It is usually used after a convolutional layer.

1	-1	-1
-1	1	-1



-1	-1	1
----	----	---


-1	1
-1	1

-1	1	-1
----	---	----

Filter 1

Filter 2





Why Pooling

- Subsampling pixels will not change the object

bird

bird

Subsampling




We can subsample the pixels to make image smaller
fewer parameters to characterize the image

- Convolutional layers in a convolutional neural network systematically apply learned filters to input images in order to create feature maps that summarize the presence of those features in the input.
- Convolutional layers prove very effective, and stacking convolutional layers in deep models allows layers close to the input to learn low-level features (e.g. lines) and layers deeper in the model to learn high-order or more abstract features, like shapes or specific objects.
- A limitation of the feature map output of convolutional layers is that they record the precise position of features in the input. This means that small movements in the position of the feature in the input image will result in a

Conv

Max
Pooling


New image but
smaller



2 x 2 image


Each filter is a
channel

The whole
CNN



Smaller than the
original image

The number of
channels is the
number of filters



Pooling



Convolution



Max

Pooling



Convolution



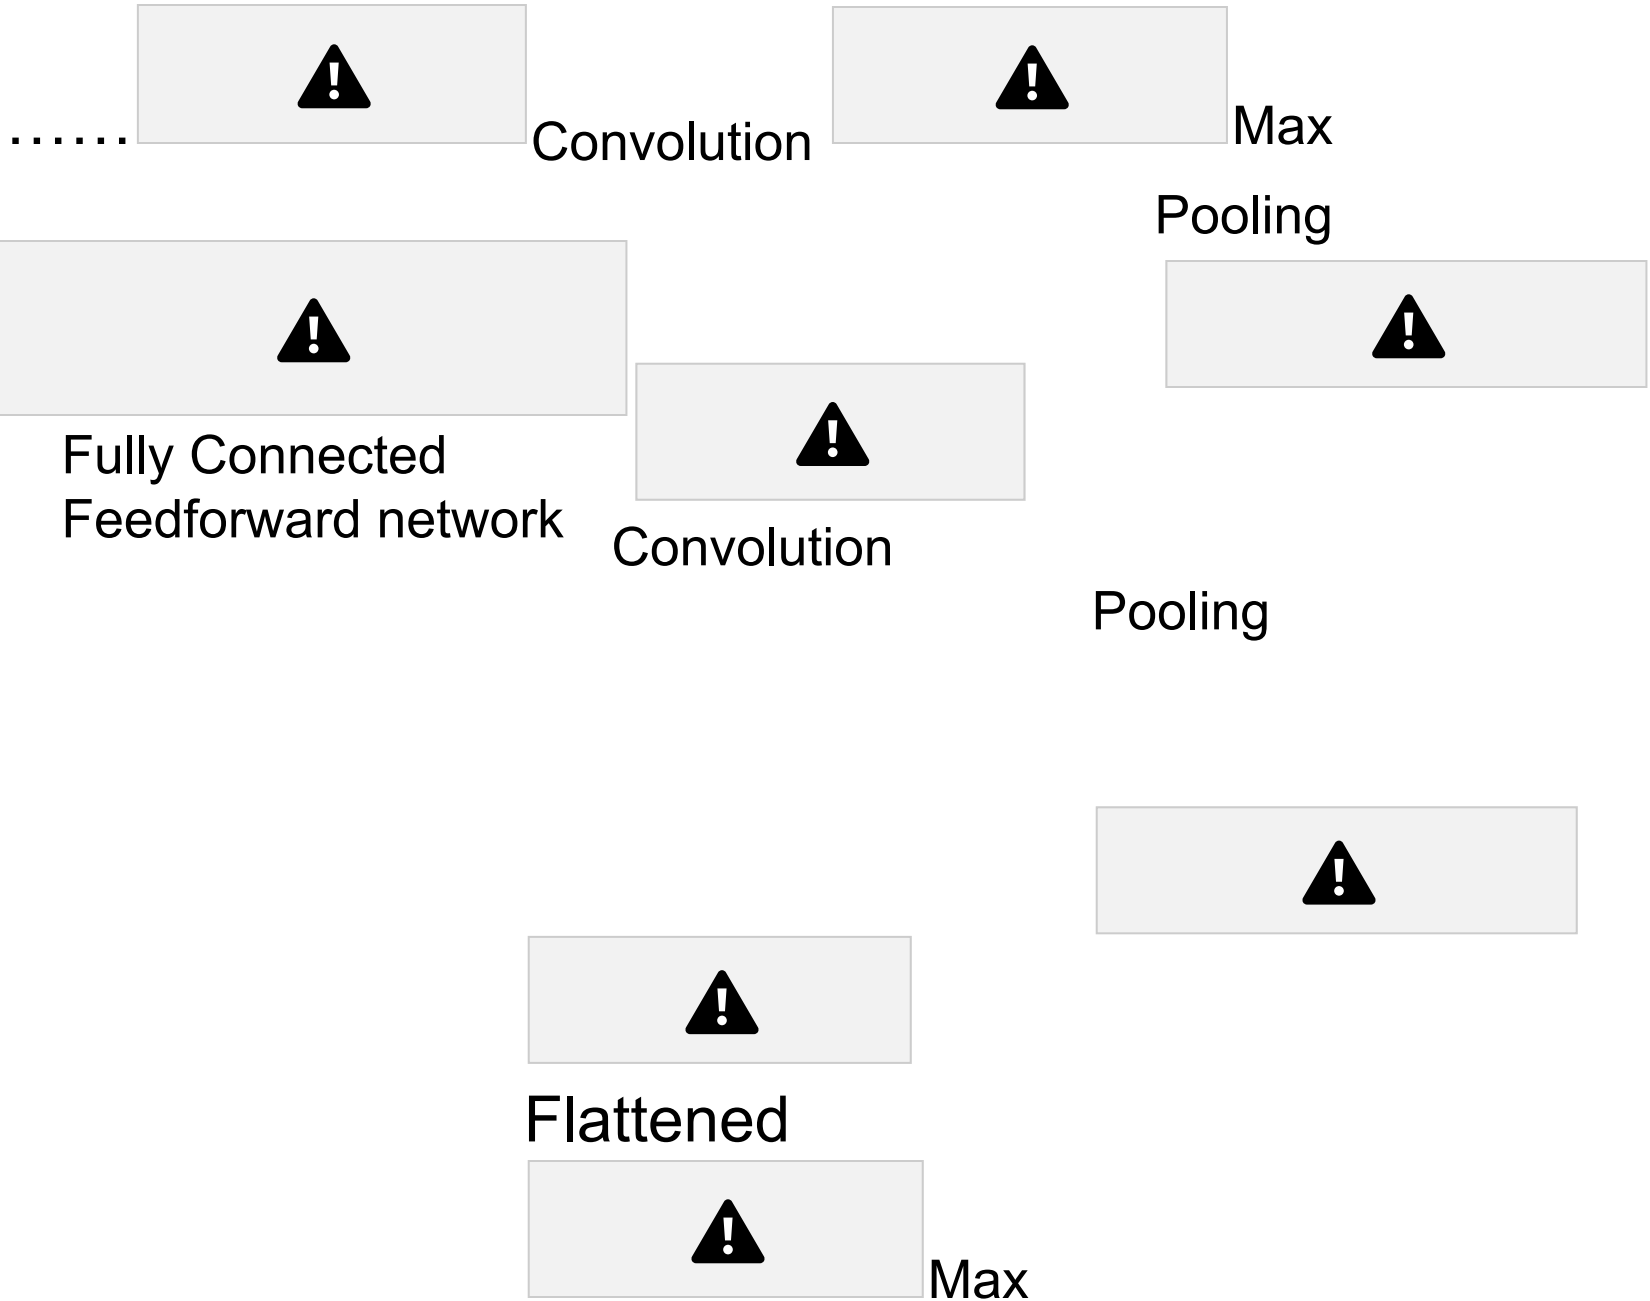
Max

Can
repeat many times


The whole CNN







cat dog



Flattening



0
3 Flattened



Fully Connected
Feedforward network



Only modified the

network


structure and

CNN in

Keras *input*

format (vector

*-> 3-D tensor)
input*



Convolution

1					
	-1	-	1	1	-1
			-		
			1		

-1	1				
		-		1	-1
-1	-1	1			
		-			
		1		1	-1

There are



...

25 3x3

...

Max Pooling


filters.

Input_shape = (28 , 28 , 1)

28 x 28 pixels 1: black/white, 3: RGB




Convolution



3 -1


3

Max Pooling



-3


1



Only modified the

network

structure and
CNN in
Keras*input*
format (vector
-> 3-D array)
 Input



x 13



Max



Convolution

Pooling

How many
parameters for
each filter?



50 x

Max Pooling



225=
25x9

50 x 11 x 11

5 x 5




Only modified
the ***network
structure*** and
CNN in

Keras *input*

format (vector

-> 3-D array)

Input

 1 x 28 x 28

Output

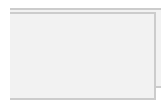


Fully connected
feedforward network





 1250



 Convolution
25 x 26 x 26




Max Pooling



AlphaGo

Neural

Next move



19 x

19 matrix

Black: 1

white: -1

none: 0

Network ^{(19 x 19} positions)



Fully-connected
feedforward network can
be used



But CNN performs much
better

AlphaGo's policy network

The following is quotation from their Nature article:



Note: AlphaGo does not use Max Pooling.






CNN in speech recognition

ye
n
e
u
n
e
r
e



CNN Image

The filters move in the frequency direction.



Time

Spectrogram

CNN in text classification



Source of image:
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.703.6858&rep=rep1&type=pdf>